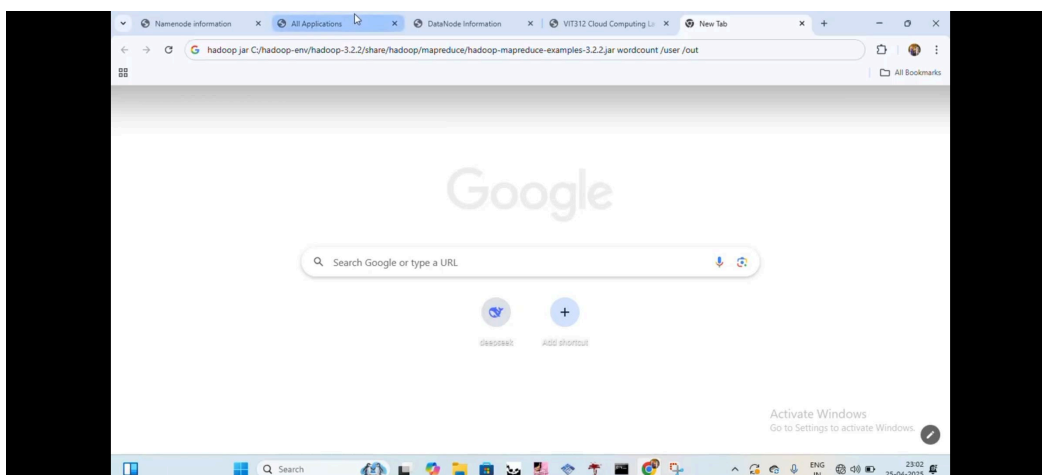


Wordcount steps

```
Administrator: Command Prompt -> C:\hadoop-env\hadoop-3.2.2\share\hadoop\mapreduce\hadoop-mapreduce-examples-3.2.2.jar wordcount /user /out
16856 DataNode
4332 Jps

C:\hadoop-env\hadoop-3.2.2\sbin>hdfs dfs -mkdir /user
C:\hadoop-env\hadoop-3.2.2\sbin>hdfs dfs -put C:/data.txt /user
C:\hadoop-env\hadoop-3.2.2\sbin>hdfs dfs -ls /user/
Found 1 items
-rw-r--r-- 1 keyan supergroup          50 2025-04-25 22:56 /user/data.txt
C:\hadoop-env\hadoop-3.2.2\sbin>hdfs dfs -cat /user/data.txt
cat: /user/data.txt: No such file or directory
C:\hadoop-env\hadoop-3.2.2\sbin>hdfs dfs -cat /user/data.txt
cloud and Grid Lab, cloud and Grid Lab, cloud Lab.
C:\hadoop-env\hadoop-3.2.2\sbin>hadoop jar C:\hadoop-env\hadoop-3.2.2\share\hadoop\mapreduce\hadoop-mapreduce-examples-3.2.2.jar wordcount /user /out
JAR does not exist or is not a normal file: C:\hadoop-env\hadoop-3.2.2\share\hadoop\mapreduce\hadoop-mapreduce-examples-3.2.2.jar
C:\hadoop-env\hadoop-3.2.2\sbin>hadoop jar C:\hadoop-env\hadoop-3.2.2\share\hadoop\mapreduce\hadoop-mapreduce-examples-3.2.2.jar wordcount /user /out
Usage: wordcount <in> [<in>...] <out>
C:\hadoop-env\hadoop-3.2.2\sbin>hadoop jar D:\hadoop-env\hadoop-3.2.2\share\hadoop\mapreduce\hadoop-mapreduce-examples-3.2.2.jar wordcount /user /out
JAR does not exist or is not a normal file: D:\hadoop-env\hadoop-3.2.2\share\hadoop\mapreduce\hadoop-mapreduce-examples-3.2.2.jar
C:\hadoop-env\hadoop-3.2.2\sbin>hadoop jar D:\hadoop-env\hadoop-3.2.2\share\hadoop\mapreduce\hadoop-mapreduce-examples-3.2.2.jar wordcount /user /out
JAR does not exist or is not a normal file: D:\hadoop-env\hadoop-3.2.2\share\hadoop\mapreduce\hadoop-mapreduce-examples-3.2.2.jar
C:\hadoop-env\hadoop-3.2.2\sbin>hadoop jar C:\hadoop-env\hadoop-3.2.2\share\hadoop\mapreduce\hadoop-mapreduce-examples-3.2.2.jar wordcount /user /out
2025-04-25 23:01:47,462 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
2025-04-25 23:01:49,267 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/keyan/.staging/job_1745601387774_0001
2025-04-25 23:01:50,257 INFO mapreduce.JobSubmitter: Total input files to process : 1
2025-04-25 23:01:50,423 INFO mapreduce.JobSubmitter: number of splits:1
2025-04-25 23:01:50,250 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1745601387774_0001
2025-04-25 23:01:50,257 INFO mapreduce.JobSubmitter: Executing with tokens: {}
2025-04-25 23:01:50,557 INFO conf.Configuration: resource-types.xml not found
2025-04-25 23:01:50,558 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2025-04-25 23:01:51,190 INFO impl.YarnClientImpl: Submitted application application_1745601387774_0001
2025-04-25 23:01:52,817 INFO mapreduce.Job: The url to track the job: http://karthi:8088/proxy/application_1745601387774_0001/
2025-04-25 23:01:52,116 INFO mapreduce.Job: Running job: job_1745601387774_0001
2025-04-25 23:02:09,693 INFO mapreduce.Job: Job job_1745601387774_0001 running in uber mode : false
2025-04-25 23:02:09,690 INFO mapreduce.Job: map 0% reduce 0%
```

```
Administrator: Command Prompt
Map input records=1
Map output records=10
Map output bytes=91
Map output materialized bytes=50
Input split bytes=100
Combine input records=10
Combine output records=4
Reduce input groups=4
Reduce shuffle bytes=50
Reduce input records=4
Reduce output records=4
Spilled Records=0
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=108
CPU time spent (ms)=1542
Physical memory (bytes) snapshot=467271608
Virtual memory (bytes) snapshot=615776256
Total committed heap usage (bytes)=308417024
Peak Map Physical memory (bytes)=299480572
Peak Map Virtual memory (bytes)=201737344
Peak Reduce Physical memory (bytes)=167723808
Peak Reduce Virtual memory (bytes)=224838912
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=50
File Output Format Counters
Bytes Written=28
C:\hadoop-env\hadoop-3.2.2\sbin>hadoop fs -cat /out/*
Grid 2
Lab 3
and 2
Cloud 3
C:\hadoop-env\hadoop-3.2.2\sbin>
```



Overview 'localhost:9820' (active)

Started:	Fri Apr 25 22:45:47 +0530 2025
Version:	3.2.2, (7a3bc50b09257c8ace2f76d74264906f07a932
Compiled:	Sun Jan 03 14:56:00 +0530 2021 by hexiaoqiao from branch-3.2.2
Cluster ID:	CID-3b92658-e6c2-472b-b2be-49ed8d51d42d
Block Pool ID:	BP-1632591371-192.168.56.1-1745601291245

Summary

Security is off.
SafeMode is off.

1 files and directories, 0 blocks (0 replicated blocks, 0 erasure coded block groups) = 1 total filesystem object(s).
Heap Memory used 94.85 MB of 236.5 MB Heap Memory. Max Heap Memory is 889 MB.
Non Heap Memory used 49.77 MB of 50.9 MB Committed Non Heap Memory. Max Non Heap Memory is <unbounded>.

Activate Windows
Go to Settings to activate Windows.

Overview 'localhost:9820' (active)

Started:	Fri Apr 25 22:45:47 +0530 2025
Version:	3.2.2, (7a3bc50b09257c8ace2f76d74264906f07a932
Compiled:	Sun Jan 03 14:56:00 +0530 2021 by hexiaoqiao from branch-3.2.2
Cluster ID:	CID-3b92658-e6c2-472b-b2be-49ed8d51d42d
Block Pool ID:	BP-1632591371-192.168.56.1-1745601291245

Summary

Security is off.
SafeMode is off.

1 files and directories, 0 blocks (0 replicated blocks, 0 erasure coded block groups) = 1 total filesystem object(s).
Heap Memory used 94.85 MB of 236.5 MB Heap Memory. Max Heap Memory is 889 MB.
Non Heap Memory used 49.77 MB of 50.9 MB Committed Non Heap Memory. Max Non Heap Memory is <unbounded>.

Activate Windows
Go to Settings to activate Windows.

- Browse the file system
- Logs
- Log Level
- Metrics
- Configuration
- Process Thread Dump

Browse Directory

Showing 1 to 3 of 3 entries

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	keyan	supergroup	0 B	Apr 25 23:02	0	0 B	log
drwxr-xr-x	keyan	supergroup	0 B	Apr 25 23:01	0	0 B	tmp
drwxr-xr-x	keyan	supergroup	0 B	Apr 25 22:56	0	0 B	user

Showing 1 to 3 of 3 entries

Previous 1 Next

Hadoop, 2021.

Activate Windows
Go to Settings to activate Windows.

